

L1: RL Basics

Deep Reinforcement Learning
Tsinghua University, 2025 Spring

Contents

1	What & RL v.s. Supervised Learning	2
2	MDP	2
2.1	Basic Concepts	2
2.2	Optimal Quantities & Bellman Equation	3
3	Value Iteration	4
4	Policy Iteration	4
4.1	Policy Evaluation	4
4.2	Policy Improvement	5
4.3	Policy Iteration	5
5	From MDP to RL	6

1 What & RL v.s. Supervised Learning

Key differences:

1. No supervision, only reward signals.
2. Rewards are delayed.
3. Not i.i.d data, but sequential.




Figure 1: RL and its common algorithms

2 MDP

2.1 Basic Concepts

Markov Decision Process (MDP) problems are non-deterministic problems.

The **definition** of MDP is as below:

- A set of states $s \in \mathcal{S}$
- A set of observations $o \in \mathcal{O}$ (sometimes)
- A set of actions $a \in \mathcal{A}$
- A transition function $T(s, a, s')$
- A reward function $R(s, a, s')$

Markov Property:

$$\begin{aligned} P(S_{t+1} = s' | S_t = s_t, A_t = a_t, \dots, S_0 = s_0) \\ = P(S_{t+1} = s' | S_t = s_t, A_t = a_t) \end{aligned} \quad (2.1)$$

Fully v.s. Partially Observability: Whether state is the same as observation. If not, transformed into POMDP.

Policy: A policy gives an action for each state. An optimal policy ($\pi^* : S \rightarrow A$) is the one that maximizes expected utility (sum of rewards).

Discount Factor: Used in calculating the sum of rewards.

Why we introduce Discount Factor:

- Avoid cyclic reward
- Infinite horizon convergence
- Worry about future uncertainties
- Time value & Human like

Utility: Sum of discounted rewards.

2.2 Optimal Quantities & Bellman Equation

The optimal policy: $\pi^*(s)$ = optimal action from a state s .

The optimal value (utility) of a state: $V^*(s)$ = expected utility starting in s and act optimally.

The optimal Q value: $Q^*(s, a)$ = expected utility taking action a from state s and acting optimally. (the action of this timestep is not necessarily the optimal)

$$\begin{aligned} V^*(s) &= \max_a Q^*(s, a) \\ Q^*(s, a) &= \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \\ V^*(s) &= \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \end{aligned} \quad (2.2)$$

3 Value Iteration

Define $V_k(s)$ to be the optimal value of s if the game ends in k more time steps, $V_0(s) = 0$, then we can compute $V_{k+1}(s)$:

$$V_k(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^k(s')] \quad (3.1)$$

Repeat this process until convergence.

NOTE

Some **problems** with value iteration:

1. Policy converges **earlier** than value.
2. The complexity of value iteration is $O(S^2 A)$.

Value iteration max over all actions to compute the optimal values. If we consider some fixed policy $\pi(s)$ (generally non-optimal), the tree would be much simpler.

Define the utility of a state s , under a fixed policy: $V_\pi(s)$ = expected total discounted rewards starting in s and following π . And the recursive relation (one-step look-ahead) is as below.

$$V_\pi(s) = \sum_{s'} T(s, \pi, s') [R(s, \pi, s') + \gamma V_\pi(s')] \quad (3.2)$$

4 Policy Iteration

4.1 Policy Evaluation

Under a fixed policy, how to calculate the V function?

Idea 1: Turn recursive Bellman equations into updates (like value itera-

tion). $O(S^2)$ complexity.

$$\begin{aligned} V_0^\pi(s) &= 0 \\ V_{k+1}^\pi(s) &= \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')] \end{aligned} \quad (4.1)$$

Idea 2: Without the maxes, the Bellman equations are just a linear system. Solve it. The complexity of matrix inversion operation (Gauss–Jordan elimination) is $O(S^3)$.

$$\begin{aligned} V &= P(r + \gamma V) \\ V &= (I - \gamma P)^{-1} R \end{aligned} \quad (4.2)$$

4.2 Policy Improvement

Imagine we have the values.

Withdraw policy from values through one-step look-up (e.g. optimal condition V^*):

$$\pi^*(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \quad (4.3)$$

One-step policy extraction:

$$\pi_{k+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')] \quad (4.4)$$

4.3 Policy Iteration

We have converged $V^{\pi_k}(s)$ through policy evaluation, then extract the next policy π_{k+1} through "argmax_a" operation traversing the entire action space. While repeating, this "a" is different from $\pi_k(s)$.

NOTE

- Policy iteration is still optimal.
- It can converge much faster under some conditions.

5 From MDP to RL

Usually we do not know the transition, thus we introduced reinforcement learning method.

Basic ideas:

1. Receive feedback in the form of rewards.
2. Utility is defined by the **reward** function.
3. Must act to maximize **expected rewards**.
4. All learning is based on observed **samples of outcomes**.

A comprehensive comparison between MDP and RL:

Aspect	MDP	RL
Definition	A mathematical model for sequential decision problems	A learning method to solve sequential decision problems
Scope	Theoretical foundation of RL; RL problems are typically modeled as MDPs	Practical application of MDP; learns to solve MDP through interaction
Known Information	Assumes transition probabilities and reward functions are known	Typically assumes transition probabilities and reward functions are unknown
Goal	Provides a framework to describe the problem	Learns the optimal policy through interaction
Methods	Based on dynamic programming (e.g., value iteration, policy iteration)	Based on trial and error (e.g., Q-learning, policy gradients)

Table 1: Comparison between MDP and RL